

SparkApplicationClass CRD proposal (Draft)

Summary

[Kubernetes Operator for Apache Spark](#) implements *SparkApplication* and *ScheduledSparkApplication* CustomResourceDefinitions (CRDs) to run or schedule spark applications in kubernetes (k8s).

SparkApplicationClass is a proposed k8s CRD to describe classes of spark applications which can be inherited by *SparkApplication* and *ScheduledSparkApplication*. This abstraction layer allows cluster administrators to declare a common set of spark application classes specific to their spark and k8s clusters, and/or their workload use cases.

API Definitions

```
ScheduledSparkApplication
|__ ScheduledSparkApplicationSpec
    |__ SparkApplication
|__ ScheduledSparkApplicationStatus
```

```
|__ SparkApplication
|__ SparkApplicationSpec
    |__ DriverSpec
        |__ SparkPodSpec
    |__ ExecutorSpec
        |__ SparkPodSpec
    |__ SparkUISpec
        |__ Dependencies
        |__ MonitoringSpec
            |__ PrometheusSpec
|__ SparkApplicationStatus
    |__ DriverInfo
```

```
SparkApplicationClass
|__ SparkApplication
```

Proposed Changes

`sparkoperator.k8s.io/class`

The updated `SparkApplication` and `ScheduledSparkApplication` will support an optional annotation `sparkoperator.k8s.io/class`. This annotation describes the name of the `SparkApplicationClass` custom resource with configurations to be inherited by the spark application.

`SparkApplication`

The updated `SparkApplication` will add 2 additional fields:

Field	Type	Note
<code>SparkApplicationClass</code>	String	Name of an existing <code>SparkApplicationClass</code> CRD that the <code>SparkApplication</code> will inherit from.
<code>SparkUI</code>	<code>SparkUISpec</code>	Optionally specifies additional properties for exposing the spark UI with k8s service and ingress.

`SparkApplicationClass`

A `SparkApplicationClass` has the following top-level field:

Field	Type	Note
Template	<code>SparkApplication</code>	A template of SparkApplication which describes the default values for a <code>SparkApplication</code> or <code>ScheduleSparkApplication</code> .

SparkUISpec

A SparkUISpec has the following top-level fields:

Field	Type	Note
Host	String	<p>Endpoint to the SparkUI (will be used to configure the Ingress endpoint).</p> <p>Supports the following mustache variables:</p> <ul style="list-style-type: none">- <code>{{appName}}</code>- <code>{{serviceName}}</code>- <code>{{namespace}}</code> <p>e.g. <code>{{appName}}.ingress.cluster.com</code></p> <p>NOTE: Available as mustache variable <code>{{host}}</code> in <code>SparkUISpec.Service</code> and <code>SparkUISpec.Ingress</code>.</p>
Service	Map[string]string	<p>An unstructured key value map which can be used by custom api gateways controllers (e.g. ambassador) to gateway to the SparkUI service.</p> <p>Supports the following mustache variables:</p> <ul style="list-style-type: none">- <code>{{appName}}</code>- <code>{{serviceName}}</code>- <code>{{namespace}}</code>- <code>{{host}}</code> <p>e.g</p> <pre>sparkUI: host: jobs.spark.com service: getambassador.io/config: --- apiVersion: ambassador/v1 kind: Mapping name: {{appName}}-{{namespace}} host: {{host}} prefix: /{{appName}}/ service: {{serviceName}}.{{namespace}}</pre>
Ingress	Map[string]string	<p>An unstructured key value map which can be used by custom ingress controllers (e.g. nginx-ingress-controller) to configure ingresses to the SparkUI service.</p>

Supports the following mustache variables:

- {{\$appName}}
- {{\$serviceName}}
- {{\$namespace}}
- {{\$host}}

e.g

sparkUI:

host: {{\$appName}}.ingress.cluster.com

ingress:

kubernetes.io/ingress.class: "nginx"

Example

SparkApplication inherits defaults from SparkApplicationClass *pyspark2*.

```
apiVersion: "sparkoperator.k8s.io/v1beta1"
kind: SparkApplication
metadata:
  name: pyspark-pi
spec:
  sparkApplicationClass: pyspark2 # inherits from SparkApplicationClass "pyspark2"
  mainApplicationFile: local:///opt/spark/examples/src/main/python/pi.py
  driver: # overwrites driver settings
    cores: "500m"
    coreLimit: "1000m"
  executor: # overwrites executor settings
    cores: 1
    instances: 20
    memory: "1Gi"
```

SparkApplicationClass with basic configuration for pyspark (with python 2) and annotations for sparkUI service and ingress.

```
apiVersion: "sparkoperator.k8s.io/v1beta1"
kind: SparkApplicationClass
metadata:
  name: pyspark2
  namespace: sparkoperator
spec:
  template:
    type: Python
    pythonVersion: "2"
    mode: cluster
    image: "gcr.io/spark-operator/spark-py:v2.4.0"
    imagePullPolicy: Always
    sparkVersion: "2.4.0"
    restartPolicy:
      type: OnFailure
      onFailureRetries: 3
      onFailureRetryInterval: 10
      onSubmissionFailureRetries: 5
      onSubmissionFailureRetryInterval: 20
  driver:
    cores: 0.1
    coreLimit: "200m"
    memory: "512m"
    labels:
```

```
    version: 2.4.0
  serviceAccount: spark
executor:
  cores: 1
  instances: 1
  memory: "512m"
  labels:
    version: 2.4.0
sparkUI: # defines default configs for sparkUI
  host: "${appName}.sparkui.com" # represented as ${host}
  service:
    getambassador.io/config: |
      ---
      apiVersion: ambassador/v1
      kind: Mapping
      name: ${appName}-${namespace}
      host: ${host}
      prefix: /
      service: ${serviceName}.${namespace}
  ingress:
    kubernetes.io/ingress.class: "nginx"
```